

CQS Summer Institute: Machine Learning and Statistics in R

Matthew S. Shotwell, Ph.D.


Department of Biostatistics
Vanderbilt University Medical Center
Nashville, TN, USA

August 14, 2018

Course Overview

- ▶ Syllabus and R code:
- ▶ <https://github.com/biostatmatt/cqs-ml-stat-r>
- ▶ Monday: Intro and Data Management
- ▶ Tuesday: Supervised Learning Part 1
- ▶ Wednesday: Supervised Learning Part 2
- ▶ Thursday: Unsupervised Learning
- ▶ Friday: Statistical Inference

Machine learning




source: <https://www.wordstream.com/blog/ws/2017/07/28/machine-learning-applications>

Supervised learning

- ▶ Have input ('features') AND output ('target')
- ▶ Create a model ('learner') using observed inputs and outputs
- ▶ Goal is to predict outputs from new inputs
- ▶ "Supervised" because both inputs *and outputs* to guide model

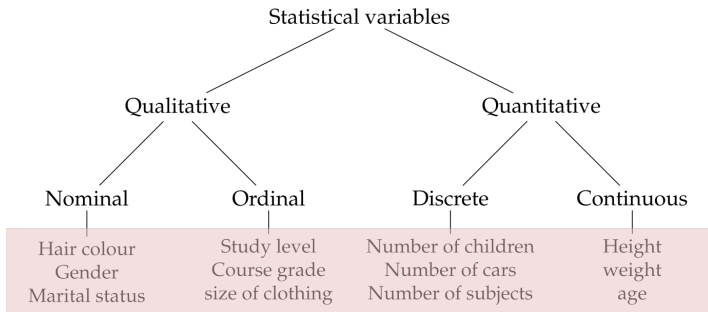
Supervised Learning Workflow



source: <https://www.quora.com/What-is-pattern-recognition>

Definitions: variable types

- ▶ quantitative - e.g. blood pressure
- ▶ qualitative - e.g. gender, a.k.a. categorical, discrete, factor, numeric codes for qualitative variables called 'targets'
- ▶ ordered - e.g. numerical pain scale (0-10)



Definitions: supervised learning tasks

- ▶ regression - model to predict quantitative output
- ▶ classification - model to predict qualitative output
- ▶ regression or classification - ordered output

Definitions: notation

- ▶ inputs - X

$$X = \begin{bmatrix} x_{11} & \cdots & x_{1p} \\ \vdots & \ddots & \vdots \\ x_{n1} & \cdots & x_{np} \end{bmatrix}$$


- ▶ quantitative outputs - Y

$$Y = \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix}$$

- ▶ qualitative outputs - G
- ▶ transpose - X^T
- ▶ prediction - \hat{Y}

Regression

- Find function f to predict Y : $\hat{Y} = \hat{f}(X)$



Regression

- ▶ To find f in $\hat{Y} = \hat{f}(X)$, need an **objective function**:
- ▶ Least squared error: minimize $1/n \sum_{i=1}^n (y_i - f(x_i))^2$
- ▶ Least absolute error: minimize $1/n \sum_{i=1}^n |y_i - f(x_i)|$
- ▶ and need to select a class of models $f(X)$:
- ▶ linear models:

$$f(X) = X\beta$$

- ▶ k-nearest neighbor:

$$f(X) = \frac{1}{k} \sum_{x_i \in N_k(X)} y_i$$

Linear regression (LS)

- Given $n \times p$ matrix X , predict Y ($n \times 1$ matrix) as follows:

$$\hat{Y} = \hat{f}(X) = X\hat{\beta}$$

where $\hat{\beta}$ is the value that minimizes the sum of squared error in the training data residual mean of squared errors:

$$1/n \sum_{i=1}^n (y_i - x_i\beta)^2$$

- The R function 'lm' will estimate β in this way

k-nearest neighbor regression (kNN)

- Predict \hat{Y} corresponding to X by averaging the Y values of the k nearest neighbors to X :

$$\hat{Y}(X) = \frac{1}{k} \sum_{x_i \in N_k(X)} y_i$$

- $N_k(X)$ is set of k points nearest X , as determined by a distance metric, e.g., the Euclidean distance:

$$d(X, X') = \sqrt{(X - X')^T (X - X')}$$

- k parameters is a 'smoothing parameter' or 'tuning parameter'
- 'caret::knnreg' implements kNN regression with Euclidean dist.

LS vs. NN method: flexibility, bias-variance

- ▶ LS is a parametric method; limited flexibility
- ▶ NN is a semiparametric method; very flexible, tunable using k
- ▶ LS cannot automatically discover complex associations
- ▶ NN can automatically discover complex associations (given enough data)
- ▶ However, LS can be made more flexible by making the linear predictor more flexible, using 1) more predictors, 2) predictor interactions, and 3) nonlinear transformations of predictors (e.g., splines)
- ▶ more flexibility results in less bias, more variance in estimator

R Code for LS and kNN

Classification (binary case)

- ▶ Think of \hat{Y} as the probability of a '1' outcome
- ▶ The predicted class \hat{G} is '1' if $\hat{Y} > 0.5$ and '0' otherwise
- ▶ Can often estimate \hat{Y} in a manner similar to regression (e.g., LS and NN).

Classification example

- ▶ Y - qualitative (binary) outcome (orange - 1, blue - 0)
- ▶ X_1 - quantitative predictor
- ▶ X_2 - quantitative predictor
- ▶ classification rule: $\hat{G} = \text{orange}$ if $\hat{Y} > 0.5$ else blue

Linear classification

Elements of Statistical Learning (2nd Ed.) ©Hastie, Tibshirani & Friedman 2009 Chap 2




FIGURE 2.1. A classification example in two dimensions. The classes are coded as a binary variable (BLUE = 0, ORANGE = 1), and then fit by linear regression. The line is the decision boundary defined by $x^T \hat{\beta} = 0.5$. The orange shaded region denotes that part of input space classified as ORANGE, while the blue region is classified as BLUE.

k-nearest neighbor classification

Elements of Statistical Learning (2nd Ed.) ©Hastie, Tibshirani & Friedman 2009 Chap 2




FIGURE 2.2. The same classification example in two dimensions as in Figure 2.1. The classes are coded as a binary variable (BLUE = 0, ORANGE = 1) and then fit by 15-nearest-neighbor averaging as in (2.8). The predicted class is hence chosen by majority vote amongst the 15-nearest neighbors.




FIGURE 2.3. The same classification example in two dimensions as in Figure 2.1. The classes are coded as a binary variable (BLUE = 0, ORANGE = 1), and then predicted by 1-nearest-neighbor classification.

Evaluating the model and ‘optimism’

- ▶ Need a mechanism to evaluate predictive quality of model and thus tune the model.
- ▶ The **prediction error** of a model is evaluated using the same objective function for estimation (e.g., mean of squared errors: $1/n \sum_{i=1}^n (y_i - x_i \beta)^2$). However, this can be an *optimistic* estimate if computed using the training data.
- ▶ The **training error** is computed on the training data
- ▶ The **test error** is computed on new data
- ▶ The **optimism** is the difference in training and test error
- ▶ Training error is optimistic, especially when the model is overfitted. It should not be used to select tuning parameters: for k-NN, $k = 1$ results in zero training error; the model is “overfitted”
- ▶ Need an estimate of test error to select tuning parameters




FIGURE 2.4. Misclassification curves for the simulation example used in Figures 2.1, 2.2 and 2.3. A single training sample of size 200 was used, and a test sample of size 10,000. The orange curves are test and the blue are training error for k -nearest-neighbor classification. The results for linear regression are the bigger orange and blue squares at three degrees of freedom. The purple line is the optimal Bayes error rate.





FIGURE 2.11. *Test and training error as a function of model complexity.*

Estimating test error: testing/training split

- ▶ Split data into training and testing data
- ▶ Use training data to build models
- ▶ Use testing data to evaluate models
- ▶ Simple, but not most efficient use of data


Estimating test error: k-fold cross validation

- ▶ Split data into k subsets or 'folds'
- ▶ Use k-1 subsets to build model
- ▶ Use holdout subset to evaluate model
- ▶ Repeat for all k permutations
- ▶ Results in k estimates of test error; use mean and sd
- ▶ More complicated, but also more efficient use of data



Estimating test error: combined CV and train/test


Supervised learning workflow can involve both train/test and CV:



Many model-building functions do cross-validation automatically.

Estimating test error: k-fold CV one-SD rule

- ▶ k-fold CV generates k estimates
- ▶ Can use their SD to gauge uncertainty
- ▶ When tuning a model, use the one-SD rule for parsimony
- ▶ Select least complex model whose CV error is within one SD of the best model



R Code for kNN 5-fold CV

Classification and regression trees (CART)

- ▶ trees partition feature space into a set of rectangles
- ▶ fit simple model in each partition (e.g., constant)
- ▶ for regression, mean within each partition
- ▶ for binary classification, compute probability within each partition
- ▶ trees are flexible like kNN, can automatically discover complexity

- ▶ top-left: partitions can't be described by tree
- ▶ top-right: partitions are recursive, can be described by tree
- ▶ bottom-left: tree describing top-right partitions
- ▶ bottom-right: regression surface for top-right partitions
- ▶ trees have (root, internal, and terminal) nodes and branches
- ▶ terminal nodes also called leaf node




FIGURE 9.2. Partitions and CART. Top right panel shows a partition of a two-dimensional feature space by recursive binary splitting, as used in CART, applied to some fake data. Top left panel shows a general partition that cannot be obtained from recursive binary splitting. Bottom left panel shows the tree corresponding to the partition in the top right panel, and a perspective plot of the prediction surface appears in the bottom right panel.

Tree-growing and tuning parameter

- ▶ step 1. find optimal split (on all possible variables)
- ▶ step 2. within each resulting partition find next optimal split
- ▶ repeat steps 1 and 2 until stopping criterion met
- ▶ tree complexity is a tuning parameter
- ▶ one approach: grow a large tree, stopping only when minimum node size (i.e., minimum number of $x_i \in R_m$) is reached, then “prune” tree back using a “cost-complexity” criterion

Pros and cons of (single) trees

Pros:

- ▶ interpretable (nice tree structure)
- ▶ flexible
- ▶ easily handle quant. and qual. data
- ▶ simple to implement

Cons:

- ▶ instability; sample variability in tree structure
- ▶ lack of smoothness of prediction surface in feature space
- ▶ categorical features difficult to split

R Code for CART: 'rpart'

Random Forests

- ▶ Random forests fits many trees to the same dataset and aggregates predictions across trees. Each tree is fit to a slightly modified version of the dataset, which is generated by resampling (bootstrap). This is generically called ‘bagging’ or ‘bootstrap aggregation’.
- ▶ regression: average predictions
- ▶ classification: average probs. or “committee vote”
- ▶ reduces variance of estimated predictor

Random Forests

- ▶ resamples are not independent samples
- ▶ bagged trees are correlated
- ▶ weakens the “wisdom of crowds”
- ▶ random forests generates *de-correlated* trees
- ▶ performance similar to boosting for many problems

Random Forests

- ▶ in bagging, trees are identically distributed (i.d.)
- ▶ → bias of bagged trees is same as individual tree
- ▶ → can only hope to improve variance
- ▶ in boosting, trees are adaptive, not i.d.

Random Forests

- ▶ variance of the average of B i.d. random variables with variance σ^2 and pairwise correlation ρ :

$$\rho\sigma^2 + \frac{1-\rho}{B}\sigma^2$$


- ▶ as B grows, second term vanishes
- ▶ can only reduce first term by reducing ρ
- ▶ random forests: reduce ρ without increasing σ^2 too much

Random Forests

- ▶ random forests modified the tree-growing procedure
- ▶ before each split, select $m \leq p$ input variables at random as candidates for splitting
- ▶ typically $m = \lfloor \sqrt{p} \rfloor$ for regression
- ▶ typically $m = \lfloor p/3 \rfloor$ for classification
- ▶ random forest predictor after B trees grown is

$$\hat{f}_{\text{rf}}^B(x) = \frac{1}{B} \sum_{b=1}^B T(x; \Theta_b)$$

Resampling (bootstrap)

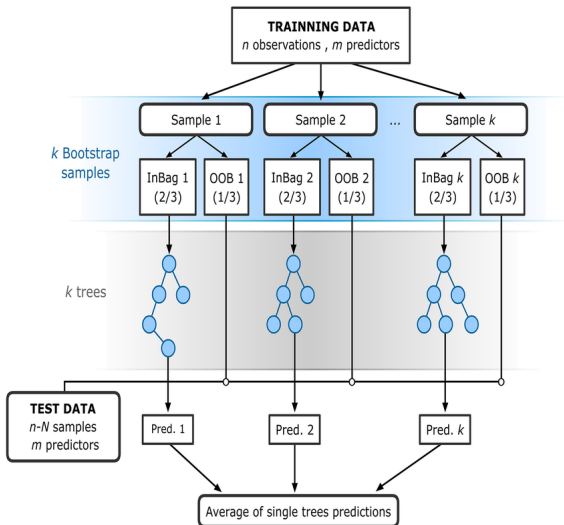


source: <https://slideplayer.com/slide/8340508/>

Random forests tree building

- ▶ before each tree grown, resample is split into training (in bag) and testing (out of bag; OOB); built-in CV
- ▶ random forests modifies the tree-growing procedure
- ▶ before each split, select $m \leq p$ input variables at random as candidates for splitting

Random forest



Pros and cons of random forest

Pros:

- ▶ more stable (less variance) than single trees
- ▶ better predictive performance
- ▶ flexible
- ▶ easily handle quant. and qual. data

Cons:

- ▶ complex to implement
- ▶ less interpretable (cannot represent as tree)
- ▶ lack of smoothness of prediction surface in feature space
- ▶ categorical features difficult to split

R Code for random forest:
`randomForest`