

# CQS Summer Institute: Machine Learning and Statistics in R

Matthew S. Shotwell, Ph.D.

Department of Biostatistics  
Vanderbilt University Medical Center  
Nashville, TN, USA

August 15, 2018

# Course Overview

- ▶ Syllabus and R code:
- ▶ <https://github.com/biostatmatt/cqs-ml-stat-r>
- ▶ Monday: Intro and Data Management
- ▶ Tuesday: Supervised Learning Part 1
- ▶ Wednesday: Supervised Learning Part 2
- ▶ Thursday: Unsupervised Learning
- ▶ Friday: Statistical Inference

# Boosting

- ▶ combines many “weak” learners  $\rightarrow$  powerful “committee”
- ▶ iteratively add “weak” learners by targeting regions of the input space where predictions were poor at previous iteration
- ▶ binary classification example: AdaBoost.M1

# AdaBoost.M1

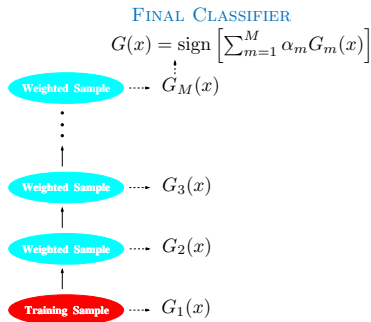
- ▶ AdaBoost.M1: popular boosted tree-based binary classifier
- ▶ binary output:  $Y \in \{-1, 1\}$
- ▶ predictors:  $X$
- ▶ classifier:  $G(X)$
- ▶ boosting is to sequentially apply a weak classifier to repeatedly modified versions of the data, thereby producing a sequence of weak classifiers  $G_m(x)$  for  $m = 1, 2, \dots, M$ .

# AdaBoost.M1

- ▶ the sequence of weak classifiers is combined using weighted majority vote:

$$G(x) = \text{sign} \left( \sum_{m=1}^M \alpha_m G_m(x) \right)$$

- ▶ weights  $\alpha_m$  are selected as part of boosting algorithm; they upweight more accurate classifiers



**FIGURE 10.1.** *Schematic of AdaBoost. Classifiers are trained on weighted versions of the dataset, and then combined to produce a final prediction.*

# AdaBoost.M1

- ▶ at each iteration, training data are reweighted
- ▶ initially weights  $w_1, \dots, w_N = 1/N$
- ▶ weak learner is then applied to weighted training data
- ▶ at iteration  $m$ , observations misclassified by  $G_{m-1}(x)$  get larger weights, and vice versa
- ▶ observations that are repeatedly misclassified get larger and larger weights
- ▶ thus, the weak learner becomes more focused on those misclassified observations

---

**Algorithm 10.1** *AdaBoost.M1*.

---

1. Initialize the observation weights  $w_i = 1/N$ ,  $i = 1, 2, \dots, N$ .
2. For  $m = 1$  to  $M$ :
  - (a) Fit a classifier  $G_m(x)$  to the training data using weights  $w_i$ .
  - (b) Compute

$$\text{err}_m = \frac{\sum_{i=1}^N w_i I(y_i \neq G_m(x_i))}{\sum_{i=1}^N w_i}.$$

- (c) Compute  $\alpha_m = \log((1 - \text{err}_m)/\text{err}_m)$ .
    - (d) Set  $w_i \leftarrow w_i \cdot \exp[\alpha_m \cdot I(y_i \neq G_m(x_i))]$ ,  $i = 1, 2, \dots, N$ .
  3. Output  $G(x) = \text{sign} \left[ \sum_{m=1}^M \alpha_m G_m(x) \right]$ .
-



# AdaBoost.M1

- ▶  $\alpha_m$  is log odds of correct classification by  $G_m(x)$
- ▶  $\text{err}_m$  always  $\leq 0.5$ , thus  $a_m \geq 0$
- ▶ weight update:

$$w_i \leftarrow w_i \exp[\alpha_m I(y_i \neq G_m(x_i))]$$


$$w_i \leftarrow \begin{cases} w_i \left( \frac{1 - \text{err}_m}{\text{err}_m} \right) & \text{if } y_i \text{ misclassified} \\ w_i & \text{otherwise} \end{cases}$$

# AdaBoost.M1 example

- ▶ let features  $X_1, \dots, X_{10}$  be standard independent normal variates
- ▶ let target  $Y$  be deterministic such that

$$y = \begin{cases} 1 & \text{if } \sum_{j=1}^{10} X_j^2 > \chi_{10}^2(0.5) \\ -1 & \text{otherwise} \end{cases}$$

- ▶ model is not additive in inputs
- ▶ high order interactions of inputs
- ▶ use “stump” as weak learner: a tree with just one split



**FIGURE 10.2.** *Simulated data (10.2): test error rate for boosting with stumps, as a function of the number of iterations. Also shown are the test error rate for a single stump, and a 244-node classification tree.*

# Gradient Boosted Models (gbm)

- ▶ `'gbm::gbm'` function
- ▶ gradient boosting for both classification and regression
- ▶ many additional options
- ▶ will automatically do k-fold CV

# AdaBoost.M1 and 'gbm' R

## Code: boosting-trees.R

# Neural Networks

A neural network is a nonlinear model, often represented using a network diagram:




FIGURE 11.2. Schematic of a single hidden layer, feed-forward neural network.

# Neural Networks

The model formula for the NN in the previous figure, in a  $K$  class configuration, is as follows:

$$Z_m = \sigma(\alpha_{0m} + \alpha_m^T X), \quad m = 1, \dots, M,$$


$$T_k = \beta_{0k} + \beta_k^T Z, \quad k = 1, \dots, K,$$

$$f_k(X) = g_k(T), \quad k = 1, \dots, K,$$

$$g_k(T) = \frac{e^{T_k}}{\sum_{\ell=1}^K e^{T_\ell}}.$$

# $\sigma$ Activation function

$\sigma$  is an “activation function” designed to mimic the behavior of neurons in propagating signals in the (human) brain.



**FIGURE 11.3.** Plot of the sigmoid function  $\sigma(v) = 1/(1 + \exp(-v))$  (red curve), commonly used in the hidden layer of a neural network. Included are  $\sigma(sv)$  for  $s = \frac{1}{2}$  (blue curve) and  $s = 10$  (purple curve). The scale parameter  $s$  controls the activation rate, and we can see that large  $s$  amounts to a hard activation at  $v = 0$ . Note that  $\sigma(s(v - v_0))$  shifts the activation threshold from 0 to  $v_0$ .

- ▶ sigmoid -  $\sigma(x) = \frac{1}{1 + e^{-x}}$
- ▶ ReLU -  $\sigma(x) = \max(0, x)$
- ▶ ReLU - faster training vs sigmoid, and may be more like real neurons



# Training (fitting) neural networks

- ▶ NN's often have large number of parameters:  $\theta$
- ▶ Regression: minimize  $R(\theta) = \sum_{i=1}^N (y_i - f(x_i, \theta))^2$
- ▶ Classification: minimize  $R(\theta) = - \sum_{i=1}^N \sum_{k=1}^K y_{ik} \log f_k(x_i, \theta)$ , where  $y_{ik}$  is the indicator for class  $k$ , and  $f_k$  gives the probability for class  $k$ . This is called the “cross-entropy” or “deviance”
- ▶  $R(\theta)$  is optimized using gradient descent algorithm called “back-propagation” or “backprop”

# Training (fitting) neural networks

- ▶ Backprop iterates two steps:
- ▶ Forward step: fix  $\theta$  and compute  $\hat{f}(x_i, \theta)$
- ▶ Backward step: fix  $\hat{f}(x_i, \theta)$  and update  $\theta$
- ▶ Using chain rule, backward step easy with gradient descent
- ▶ Usually don't want global minimum of  $R(\theta)$  due to overfitting
- ▶ # of iters, learning rate, stopping criteria, and shrinkage (weight decay, dropout) are tuning parameters
- ▶ Can be modified for big data and parallelization
- ▶ Backprop is an “art” (starting values for  $\theta$ , regularization using “weight decay”, etc.)


# Simple NN in R: nnet.R

# Extending NN's

- ▶ The real power of NN's comes through various extensions:
- ▶ Additional hidden layers
- ▶ Modifying connectivity between layers
- ▶ Processing between layers

# Additional layers

More than one hidden layer:



# Modified connectivity

- ▶ Local connectivity: hidden units do not receive input from all units in the layer below; not “fully” connected
- ▶ Weight sharing: some hidden units weight their inputs the same way, e.g., for hidden unit  $j$  and  $k$ ,  $\alpha_j = \alpha_k$ ; they share weights:


$$Z_m = \sigma(\alpha_{0m} + \alpha_m^T X), \quad m = 1, \dots, M,$$

$$T_k = \beta_{0k} + \beta_k^T Z, \quad k = 1, \dots, K,$$


$$f_k(X) = g_k(T), \quad k = 1, \dots, K,$$

## Example: zipcode data

- ▶ hand-written integers
- ▶ output: 10-class classification
- ▶ input: 16x16 B&W image




# Local connectivity






# Local connectivity and shared weights

AKA: convolutional neural networks



- ▶ groups of hidden units form “shape detectors”
- ▶ more complex shape detectors near output layer


# Performance on zipcode data



**FIGURE 11.11.** Test performance curves, as a function of the number of training epochs, for the five networks of Table 11.1 applied to the ZIP code data. (Le Cun, 1989)


# Processing between layers

- Pooling/subsampling: down-sample data from a layer by summarizing of a group of units
- Max-pooling: summarize using maximum:



# Deep learning and deep neural networks

- ▶ Deep learning uses deep NNs
- ▶ Deep NNs are simply NNs with many layers, complex connectivity, and processing steps between layers:



# Complex NNs in R

- ▶ No (good) native R libraries for complex NNs
- ▶ R can interface to good libraries, e.g., TensorFlow
- ▶ See <https://tensorflow.rstudio.com/>
- ▶ Deep NNs are simply NNs with many layers, complex connectivity, and processing steps between layers:

# Deep dream

- ▶ Start with a trained deep NN designed to classify images, e.g., of cats versus dogs
- ▶ Input a new image and evaluate all of the hidden and output layers
- ▶ Use a backprop-like process to update the hidden layers and the input image (rather than the model weights)
- ▶ The input image will be updated in such a way that the model prediction is more confident, i.e., to look more like an example from one of the classes (e.g., a cat or dog)

# Deep dream

Jellyfish picture processed using deep NN trained to distinguish breeds of dogs:



# Combining data

Paper: <https://arxiv.org/abs/1508.06576>





**TensorFlow in R:** <https://tensorflow.rstudio.com/keras/>